

## 2.1. Problem Statement: Machine Learning 7

In this assignment, students will be using the K-nearest neighbors algorithm to predict how many points NBA players scored in the 2013-2014 season.

A look at the data Before we dive into the algorithm, let's take a look at our data. Each row in the data contains information on how a player performed in the 2013-2014 NBA season.

**Here are some selected columns from the data:**

- player - name of the player pos - the position of the player g - number of games the player was in gs - number of games the player started pts - total points the player scored

**There are many more columns in the data, mostly containing information about average player game performance over the course of the season.**

See this site for an explanation of the rest of them. We can read our dataset in and figure out which columns are present:

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsRegressor
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
import math
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: nba = pd.read_csv('nba_2013.csv')
```

```
In [3]: nba.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 481 entries, 0 to 480
Data columns (total 31 columns):
#   Column          Non-Null Count  Dtype
---  -
0   player          481 non-null   object
1   pos             481 non-null   object
2   age             481 non-null   int64
3   bref_team_id    481 non-null   object
4   g               481 non-null   int64
5   gs              481 non-null   int64
6   mp              481 non-null   int64
7   fg              481 non-null   int64
8   fga             481 non-null   int64
9   fg.             479 non-null   float64
10  x3p             481 non-null   int64
```

```

11  x3pa          481 non-null    int64
12  x3p.          414 non-null    float64
13  x2p           481 non-null    int64
14  x2pa          481 non-null    int64
15  x2p.          478 non-null    float64
16  efg.          479 non-null    float64
17  ft            481 non-null    int64
18  fta           481 non-null    int64
19  ft.           461 non-null    float64
20  orb           481 non-null    int64
21  drb           481 non-null    int64
22  trb           481 non-null    int64
23  ast           481 non-null    int64
24  stl           481 non-null    int64
25  blk           481 non-null    int64
26  tov           481 non-null    int64
27  pf            481 non-null    int64
28  pts           481 non-null    int64
29  season        481 non-null    object
30  season_end    481 non-null    int64
dtypes: float64(5), int64(22), object(4)
memory usage: 116.6+ KB


```

In [4]: `nba.describe()`

Out[4]:

	age	g	gs	mp	fg	fga	fg.	x3p
<b>count</b>	481.000000	481.000000	481.000000	481.000000	481.000000	481.000000	479.000000	481.000000
<b>mean</b>	26.509356	53.253638	25.571726	1237.386694	192.881497	424.463617	0.436436	39.613306
<b>std</b>	4.198265	25.322711	29.658465	897.258840	171.832793	368.850833	0.098672	50.855639
<b>min</b>	19.000000	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	23.000000	32.000000	0.000000	388.000000	47.000000	110.000000	0.400500	0.000000
<b>50%</b>	26.000000	61.000000	10.000000	1141.000000	146.000000	332.000000	0.438000	16.000000
<b>75%</b>	29.000000	76.000000	54.000000	2016.000000	307.000000	672.000000	0.479500	68.000000
<b>max</b>	39.000000	83.000000	82.000000	3122.000000	849.000000	1688.000000	1.000000	261.000000

8 rows × 27 columns



In [5]: `nba.columns`

Out[5]: Index(['player', 'pos', 'age', 'bref\_team\_id', 'g', 'gs', 'mp', 'fg', 'fga', 'fg.', 'x3p', 'x3pa', 'x3p.', 'x2p', 'x2pa', 'x2p.', 'efg.', 'ft', 'fta', 'ft.', 'orb', 'drb', 'trb', 'ast', 'stl', 'blk', 'tov', 'pf', 'pts', 'season', 'season\_end'], dtype='object')

In [6]: `nba.isnull().sum()`

Out[6]:

player	0
pos	0
age	0
bref_team_id	0
g	0
gs	0
mp	0

```
fg          0
fga         0
fg.         2
x3p         0
x3pa        0
x3p.       67
x2p         0
x2pa        0
x2p.        3
efg.        2
ft          0
fta         0
ft.        20
orb         0
drb         0
trb         0
ast         0
stl         0
blk         0
tov         0
pf          0
pts         0
season      0
season_end  0
dtype: int64
```

```
In [7]: nba[nba["x3p."].isnull()].head()
```

Out[7]:

	player	pos	age	bref_team_id	g	gs	mp	fg	fga	fg.	...	drb	trb	ast	stl	blk	tov
1	Steven Adams	C	20	OKC	81	20	1197	93	185	0.503	...	190	332	43	40	57	71
2	Jeff Adrien	PF	27	TOT	53	12	961	143	275	0.520	...	204	306	38	24	36	39
5	Cole Aldrich	C	25	NYK	46	2	330	33	61	0.541	...	92	129	14	8	30	18
11	Louis Amundson	PF	31	TOT	19	0	185	16	32	0.500	...	27	55	6	9	11	14
18	Joel Anthony	C	31	TOT	33	0	186	12	32	0.375	...	23	38	2	3	12	3

5 rows × 31 columns



```
In [8]: nba.season.value_counts()
```

Out[8]: 2013-2014 481  
Name: season, dtype: int64

```
In [9]: nba.season_end.value_counts()
```

Out[9]: 2013 481  
Name: season\_end, dtype: int64

```
In [10]: feat_list_obj = []
         for i in nba.columns:
             if nba[i].dtype=="object":
                 feat_list_obj.append(i)
```

```
feat_list_obj
```

```
Out[10]: ['player', 'pos', 'bref_team_id', 'season']
```

```
In [11]: feat_drop_list = feat_list_obj + ["season_end", "pts"]
         feat_drop_list
```

```
Out[11]: ['player', 'pos', 'bref_team_id', 'season', 'season_end', 'pts']
```

```
In [12]: Features = nba.drop(feat_drop_list,axis=1)
         Lables= nba["pts"]
         print(Features.shape)
         print(Lables.shape)
```

```
(481, 25)
```

```
(481,)
```

```
In [13]: Features.isnull().sum()
```

```
Out[13]: age      0
         g        0
         gs       0
         mp       0
         fg       0
         fga      0
         fg.      2
         x3p      0
         x3pa     0
         x3p.    67
         x2p      0
         x2pa     0
         x2p.     3
         efg.     2
         ft       0
         fta      0
         ft.    20
         orb      0
         drb      0
         trb      0
         ast      0
         stl      0
         blk      0
         tov      0
         pf       0
         dtype: int64
```

```
In [14]: Features["fg."].fillna(Features["fg."].mean(),inplace=True)
         Features["x2p."].fillna(Features["x2p."].mean(),inplace=True)
         Features["efg."].fillna(Features["efg."].mean(),inplace=True)
         Features["x3p."].fillna(Features["x3p."].mean(),inplace=True)
         Features["ft."].fillna(Features["ft."].mean(),inplace=True)
         Y = Lables.values ##--> Dependent Values
```

```
In [15]: Features.isnull().sum()
```

```
Out[15]: age      0
         g        0
         gs       0
         mp       0
         fg       0
         fga      0
```

```

fg.      0
x3p      0
x3pa     0
x3p.     0
x2p      0
x2pa     0
x2p.     0
efg.     0
ft       0
fta      0
ft.      0
orb      0
drb      0
trb      0
ast      0
stl      0
blk      0
tov      0
pf       0
dtype: int64

```

```
In [16]: x_train,x_test,y_train,y_test = train_test_split(Features,Y)
```

```
In [17]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```

(360, 25)
(121, 25)
(360,)
(121,)

```

```
In [18]: sns.set_style("whitegrid")
for k in range(20):
    k_value=1+k
    neu=igh = KNeighborsRegressor(n_neighbors=k_value)
    neu.fit(x_train,y_train)
    y_pred = neu.predict(x_test)
    print("RMSE is :- ",np.sqrt(mean_squared_error(y_pred,y_test)),"For k value {}".for
```

```

RMSE is :- 88.86125316497392 For k value 1
RMSE is :- 82.83955916096772 For k value 2
RMSE is :- 78.93953353096481 For k value 3
RMSE is :- 72.22560284699306 For k value 4
RMSE is :- 73.88683551763411 For k value 5
RMSE is :- 72.13871063502926 For k value 6
RMSE is :- 72.57319485807122 For k value 7
RMSE is :- 76.62458458490929 For k value 8
RMSE is :- 76.37109751909952 For k value 9
RMSE is :- 77.96782235200652 For k value 10
RMSE is :- 75.8551453334614 For k value 11
RMSE is :- 77.2244888054174 For k value 12
RMSE is :- 76.61007712476186 For k value 13
RMSE is :- 77.53038430710822 For k value 14
RMSE is :- 78.04177850305048 For k value 15
RMSE is :- 80.2516382730359 For k value 16
RMSE is :- 81.28241344715694 For k value 17
RMSE is :- 83.62974282094571 For k value 18
RMSE is :- 84.2171561113578 For k value 19
RMSE is :- 83.95040288544241 For k value 20

```

```
In [19]: #note : It shows that we are get less error for values K = 5 , 6
k_value=5
```

```

neigh=igh = KNeighborsRegressor(n_neighbors=k_value)
neigh.fit(x_train,y_train)
y_pred = neigh.predict(x_test)
print("RMSE for KNN Regressor is :- ",np.sqrt(mean_squared_error(y_pred,y_test)),"For
print("R Squared for KNN Regressor is :- ",r2_score(y_test,y_pred))

```

RMSE for KNN Regressor is :- 73.88683551763411 For k value 5  
 R Squared for KNN Regressor is :- 0.9703610383923194


In [20]: *#R squared is a statistical measure of how close the data points are to thr fitted regr*

```

plt.figure(figsize=(10,8))
plt.scatter(y_test,y_pred)
plt.plot([0,2500],[0,2500], '--k')
plt.axis("tight")
plt.xlabel("Actual Rating")
plt.ylabel("Predicted Rating")
plt.tight_layout()
plt.title("Actual vs Predicted Points")

```

Out[20]: Text(0.5, 1.0, 'Actual vs Predicted Points')



```

In [21]: from sklearn.ensemble import RandomForestRegressor
Rf = RandomForestRegressor(random_state = 1)
Rf.fit(x_train,y_train)
y_pred=Rf.predict(x_test)

```

```
print("RMSE for RandomForest Regressor is :- ",np.sqrt(mean_squared_error(y_pred,y_test)))
print("R Squared for RandomForest Regressor is :- ",r2_score(y_test,y_pred))
```


RMSE for RandomForest Regressor is :- 43.25424283397521

R Squared for RandomForest Regressor is :- 0.9898425129090567

In [22]: *#R squared is a statistical measure of how close the data points are to the fitted regression line*

```
plt.figure(figsize=(10,8))
plt.scatter(y_test,y_pred)
plt.plot([0,2500],[0,2500], '--k')
plt.axis("tight")
plt.xlabel("Actual Rating")
plt.ylabel("Predicted Rating")
plt.tight_layout()
plt.title("Actual vs Predicted Points--RandomForestRegressor",fontsize=18)
```

Out[22]: Text(0.5, 1.0, 'Actual vs Predicted Points--RandomForestRegressor')



```
In [23]: for depth in range(30):
          depth+=1
          Rf = RandomForestRegressor(max_depth=depth,random_state =1)
          Rf.fit(x_train,y_train)
          y_pred=Rf.predict(x_test)
          print("RMSE is :- ",np.sqrt(mean_squared_error(y_pred,y_test)),"For Depth value :-")
```

```

RMSE is :- 230.5489415724857 For Depth value :- 1
RMSE is :- 107.58148434819346 For Depth value :- 2
RMSE is :- 67.55834208828297 For Depth value :- 3
RMSE is :- 54.60281636989371 For Depth value :- 4
RMSE is :- 48.56568661113202 For Depth value :- 5
RMSE is :- 46.33814255372531 For Depth value :- 6
RMSE is :- 45.82865164515156 For Depth value :- 7
RMSE is :- 46.0572345946096 For Depth value :- 8
RMSE is :- 44.3913087385565 For Depth value :- 9
RMSE is :- 45.03475499611181 For Depth value :- 10
RMSE is :- 45.52423691405091 For Depth value :- 11
RMSE is :- 43.94553264841992 For Depth value :- 12
RMSE is :- 43.308874734690626 For Depth value :- 13
RMSE is :- 43.25424283397521 For Depth value :- 14
RMSE is :- 43.25424283397521 For Depth value :- 15
RMSE is :- 43.25424283397521 For Depth value :- 16
RMSE is :- 43.25424283397521 For Depth value :- 17
RMSE is :- 43.25424283397521 For Depth value :- 18
RMSE is :- 43.25424283397521 For Depth value :- 19
RMSE is :- 43.25424283397521 For Depth value :- 20
RMSE is :- 43.25424283397521 For Depth value :- 21
RMSE is :- 43.25424283397521 For Depth value :- 22
RMSE is :- 43.25424283397521 For Depth value :- 23
RMSE is :- 43.25424283397521 For Depth value :- 24
RMSE is :- 43.25424283397521 For Depth value :- 25
RMSE is :- 43.25424283397521 For Depth value :- 26
RMSE is :- 43.25424283397521 For Depth value :- 27
RMSE is :- 43.25424283397521 For Depth value :- 28
RMSE is :- 43.25424283397521 For Depth value :- 29
RMSE is :- 43.25424283397521 For Depth value :- 30

```

```

In [49]: #NOTE :-the random forest regressor gives a low RMSE value for maxdepth =13 and The RMS

#CONCLUSION

#The R Squared for KNN Regressor is 0.9703610383923194

#The R Squared for RandomForest Regressor is 0.9898425129090567

```

```
In [ ]:
```

```
In [ ]:
```

## 2.2. Problem Statement: Machine Learning 8

In this assignment students have to find the frequency of words in a webpage. User can use urllib and BeautifulSoup to extract text from webpage.

**Hint:**

```
from bs4 import BeautifulSoup
```

```
import urllib.request
```

```
import nltk
```



```
response = urllib.request.urlopen('http://php.net/')
```

```
html = response.read()
```

```
soup = BeautifulSoup(html,"html5lib")
```

```
In [24]: from bs4 import BeautifulSoup
import urllib.request
import nltk
```

```
In [25]: response = urllib.request.urlopen('http://php.net/')
html = response.read()
raw = BeautifulSoup(html,"html5lib").get_text()
```

```
In [26]: nltk.download('punkt')
words=nltk.word_tokenize(raw)

#removing the singal characters mostly puncatuations
words=[word for word in words if len(word)>1]

#removing any numbers present in our text
words = [word for word in words if not word.isnumeric()]

#lowercase all words (default stopwords are lowercase too)
words = [word.lower() for word in words]

#calculating frequency distribution
fdist = nltk.FreqDist(words)

#printing the top 10 words with their frequency
for word , frequency in fdist.most_common(10):
    print(u'{}; {}'.format(word , frequency))
```

```
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\idofa\AppData\Roaming\nltk_data...
the; 245
php; 153
of; 87
release; 85
for; 81
this; 66
in; 60
is; 56
to; 51
be; 50
[nltk_data] Package punkt is already up-to-date!
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

## 2.3. Problem Statement: Machine Learning 9

*In this assignment students have to compress racoon grey scale image into 5 clusters. In the end, visualize both raw and compressed image and look for quality difference.*

The raw image is available in scipy.misc package with the name face. Hint:

```
import numpy as np
```

```
from sklearn import cluster, datasets
```

```
from scipy import misc
```

```
In [27]: #Importing Libraries
import numpy as np
from sklearn.cluster import KMeans
from sklearn import datasets
from scipy import misc
import matplotlib.pyplot as plt
```

```
In [28]: face =misc.face(gray=True)
```

```
In [29]: k_Means= KMeans(n_clusters = 5)
np.random.seed(10)
X = face.reshape((-1,1))

#fitting value of X
k_Means.fit(X)

values = k_Means.cluster_centers_
print("Values :-",values)

labels = k_Means.labels_
print("Labels :-",labels)
```

```
Values :- [[114.99362851]
 [ 27.62031146]
 [194.13840989]
 [ 75.41095451]
 [153.31393344]]
Labels :- [0 0 4 ... 4 4 4]
```

```
In [30]: #Create an array from label and values
face_compressed = np.choose(labels , values)
face_compressed.shape = face.shape

vmin = face.min()
vmax = face.max()

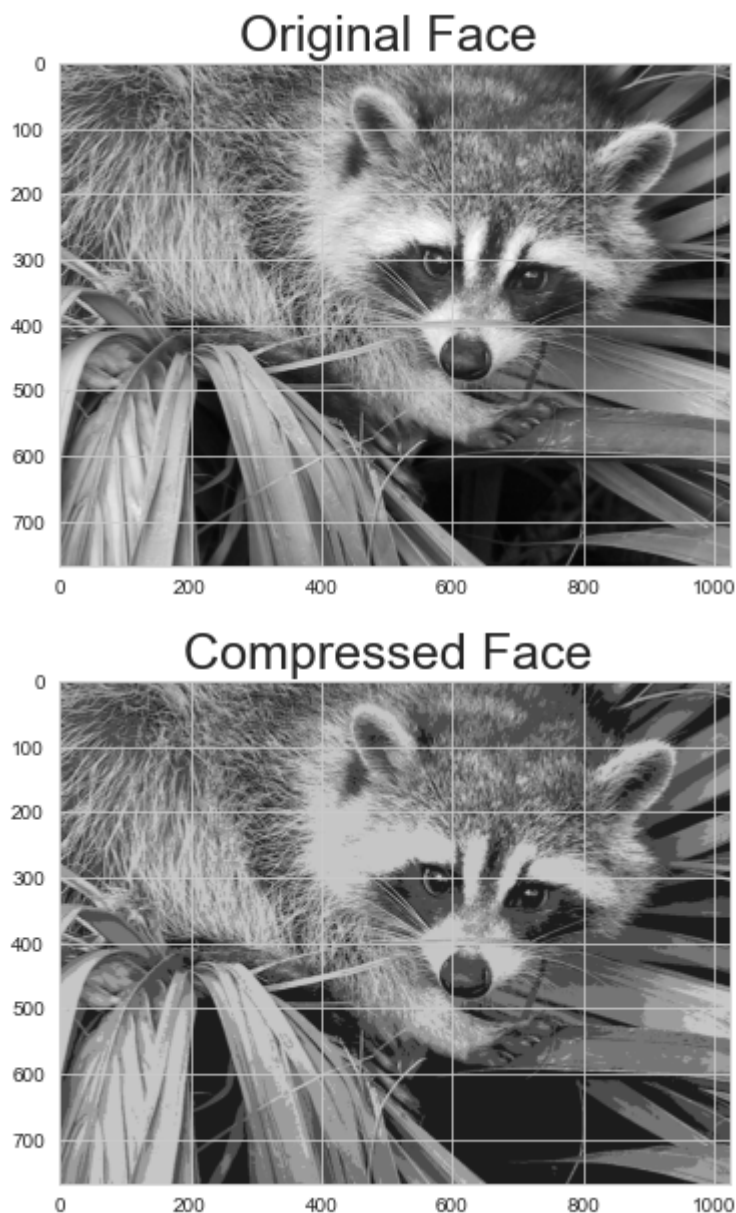
print("Vmin :-{}\t Vmax :-{}".format(vmin , vmax))
```

```
Vmin :-0          Vmax :-250
```

```
In [31]: #Original Face
plt.figure(figsize = (6,6))
plt.imshow(face,cmap= plt.cm.gray)
plt.title("Original Face ",fontsize = 25)
```

```
#Compressed Face
plt.figure(figsize = (6,6))
plt.imshow(face_compressed,cmap= plt.cm.gray,vmin = vmin,vmax= vmax)
plt.title("Compressed Face ",fontsize = 25)
```

Out[31]: Text(0.5, 1.0, 'Compressed Face ')



In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

## 2.4. Problem Statement: Machine Learning 10

***In this assignment students have to transform iris data into 3 dimensions and plot a 3d chart with transformed dimensions and colour each data point with specific class.***

Hint:

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from mpl_toolkits.mplot3d import Axes3D
```

```
from sklearn import decomposition
```

```
from sklearn import datasets
```

```
In [32]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn import decomposition
from sklearn import datasets
import seaborn as sns
import pandas as pd
```

```
In [33]: iris = sns.load_dataset('iris')
```

```
In [34]: iris.head(3)
```

```
Out[34]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
<b>0</b>	5.1	3.5	1.4	0.2	setosa
<b>1</b>	4.9	3.0	1.4	0.2	setosa
<b>2</b>	4.7	3.2	1.3	0.2	setosa

```
In [35]: iris['Species'] = pd.factorize(iris.species)[0]
```

```
In [36]: iris.tail(3)
```


```
Out[36]:
```

	sepal_length	sepal_width	petal_length	petal_width	species	Species
<b>147</b>	6.5	3.0	5.2	2.0	virginica	2
<b>148</b>	6.2	3.4	5.4	2.3	virginica	2
<b>149</b>	5.9	3.0	5.1	1.8	virginica	2

```
In [37]: plt.figure(figsize = (12,6))
ax = plt.axes(projection='3d')
ax.scatter3D(iris.sepal_length,iris.sepal_width, iris.petal_length,c =iris['Species'] ,
ax.set_xlabel('sepal_length' , fontsize = 12.5)
ax.set_ylabel("sepal_width", fontsize = 12.5)
ax.set_zlabel('petal_length', fontsize = 12.5)

Text(0.5, 0, 'petal_length')
```

Out[37]:



In [ ]:

## 2.5. Problem Statement: Machine Learning 11

***In this assignment students have to make ARIMA model over shampoo sales data and check the MSE between predicted and actual value.***

Student can download data in .csv format from the following link:

<https://datamarket.com/data/set/22r0/sales-of-shampoo-over-a-three-yearperiod#!ds=22r0&display=line>

Hint:

Following is the command import packages

and data from pandas import read\_csv

from pandas import datetime

from matplotlib import pyplot

from statsmodels.tsa.arima\_model

import ARIMA from sklearn.metrics

import mean\_squared\_error def

parser(x):

```
return datetime.strptime('190'+x, '%Y-%m')
```

```
series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True,
date_parser=parser)
```

```
In [38]: import pandas as pd
from datetime import datetime
from matplotlib import pyplot
from statsmodels.tsa.arima_model import ARIMA
from sklearn.metrics import mean_squared_error
```

```
In [39]: def parser(x):
return datetime.strptime('190'+x, '%Y-%m')
```

```
In [40]: series = pd.read_csv("sales-of-shampoo-over-a-three-ye.csv", header=0, parse_dates=True
series.head()
```


Out[40]: **Sales of shampoo over a three year period**

Month	
1-01	266.0
1-02	145.9
1-03	183.1
1-04	119.3
1-05	180.3

```
In [41]: import warnings
warnings.filterwarnings('ignore')
```

```
In [42]: series.dropna(axis=0,inplace=True)
```

```
In [43]: series.plot()
pyplot.show()
```



```
In [44]: X= series.values
X
```

```
Out[44]: array([[266. ],
               [145.9],
               [183.1],
               [119.3],
               [180.3],
               [168.5],
               [231.8],
               [224.5],
               [192.8],
               [122.9],
               [336.5],
               [185.9],
               [194.3],
               [149.5],
               [210.1],
               [273.3],
               [191.4],
               [287. ],
               [226. ],
               [303.6],
               [289.9],
               [421.6],
               [264.5],
               [342.3],
               [339.7],
               [440.4],
               [315.9],
               [439.3],
               [401.3],
               [437.4],
               [575.5],
               [407.6],
               [682. ],
               [475.3],
               [581.3],
               [646.9]])
```

```
In [45]: size = int(len(X)*0.6)
         print(len(X),'\n',size)
```

```
36
21
```

```
In [46]: train,test=X[0:size],X[size:len(X)]
```

```
In [47]: history=[x for x in train]
         predictions=[]
         for t in range(len(test)):
             model =ARIMA(history,order=(5,1,0))
             model_fit = model.fit(disp=0)
             output = model_fit.forecast()
             yhat = output[0]
             predictions.append(yhat)
             obs=test[t]
             history.append(obs)
             print("Predicted = {}    Excepted = {}".format(yhat,obs))


         error = mean_squared_error(test,predictions)
         print("\n\nTest MSE = {}".format(round(error,ndigits=3)))
```

```
Predicted = [272.96451422]    Excepted = [421.6]
Predicted = [290.31366212]    Excepted = [264.5]
Predicted = [349.11764194]    Excepted = [342.3]
```

Predicted = [306.51293429]	Excepted = [339.7]
Predicted = [387.37635503]	Excepted = [440.4]
Predicted = [348.15422407]	Excepted = [315.9]
Predicted = [386.30873564]	Excepted = [439.3]
Predicted = [356.08213522]	Excepted = [401.3]
Predicted = [446.37949922]	Excepted = [437.4]
Predicted = [394.73731958]	Excepted = [575.5]
Predicted = [434.9155186]	Excepted = [407.6]
Predicted = [507.92333692]	Excepted = [682.]
Predicted = [435.48297593]	Excepted = [475.3]
Predicted = [652.74379393]	Excepted = [581.3]
Predicted = [546.34341033]	Excepted = [646.9]

Test MSE = 7547.808

```
In [48]: pyplot.plot(test)
pyplot.plot(predictions,color='red')
pyplot.show()
```



In [ ]: