

```
In [1]: import sqlite3
import numpy as np
import pandas as pd
%matplotlib notebook
import matplotlib.pyplot as plt
import xgboost as xgb
from xgboost.sklearn import XGBRegressor
from xgboost import plot_importance

from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import Imputer, StandardScaler
from sklearn.feature_selection import SelectFromModel
from sklearn.model_selection import train_test_split, GridSearchCV, ShuffledSplit, RandomizedSearchCV
from sklearn.pipeline import make_pipeline

import pickle
```

Reading Data from the Database into pandas

```
In [2]: cnx = sqlite3.connect('database.sqlite')
```

```
In [3]: df= pd.read_sql_query("SELECT * FROM Player_Attributes", cnx)
```

```
In [4]: print(df)
```

	id	player_fifa_api_id	player_api_id	date
\				
0	1	218353	505942	2016-02-18 00:00:00
1	2	218353	505942	2015-11-19 00:00:00
2	3	218353	505942	2015-09-21 00:00:00
3	4	218353	505942	2015-03-20 00:00:00
4	5	218353	505942	2007-02-22 00:00:00
...
183973	183974	102359	39902	2009-08-30 00:00:00
183974	183975	102359	39902	2009-02-22 00:00:00
183975	183976	102359	39902	2008-08-30 00:00:00
183976	183977	102359	39902	2007-08-30 00:00:00
183977	183978	102359	39902	2007-02-22 00:00:00

	overall_rating	potential	preferred_foot	attacking_work_rate	\
0	67.0	71.0	right	medium	
1	67.0	71.0	right	medium	
2	62.0	66.0	right	medium	
3	61.0	65.0	right	medium	
4	61.0	65.0	right	medium	
...	
183973	83.0	85.0	right	medium	
183974	78.0	80.0	right	medium	
183975	77.0	80.0	right	medium	
183976	78.0	81.0	right	medium	
183977	80.0	81.0	right	medium	

```

      defensive_work_rate  crossing  ...  vision  penalties  marking
\
0      medium      49.0  ...      54.0      48.0      65.0
1      medium      49.0  ...      54.0      48.0      65.0
2      medium      49.0  ...      54.0      48.0      65.0
3      medium      48.0  ...      53.0      47.0      62.0
4      medium      48.0  ...      53.0      47.0      62.0
...      ...      ...  ...      ...      ...      ...
183973      low      84.0  ...      88.0      83.0      22.0
183974      low      74.0  ...      88.0      70.0      32.0
183975      low      74.0  ...      88.0      70.0      32.0
183976      low      74.0  ...      88.0      53.0      28.0
183977      low      74.0  ...      88.0      53.0      38.0

      standing_tackle  sliding_tackle  gk_diving  gk_handling  gk_kic
king \
0      69.0      69.0      6.0      11.0
10.0
1      69.0      69.0      6.0      11.0
10.0
2      66.0      69.0      6.0      11.0
10.0
3      63.0      66.0      5.0      10.0
9.0
4      63.0      66.0      5.0      10.0
9.0
...      ...      ...      ...      ...
...
183973      31.0      30.0      9.0      20.0
84.0
183974      31.0      30.0      9.0      20.0
73.0
183975      31.0      30.0      9.0      20.0
73.0
183976      32.0      30.0      9.0      20.0
73.0
183977      32.0      30.0      9.0      9.0
78.0

      gk_positioning  gk_reflexes
0      8.0      8.0
1      8.0      8.0
2      8.0      8.0
3      7.0      7.0
4      7.0      7.0
...      ...      ...
183973      20.0      20.0
183974      20.0      20.0
183975      20.0      20.0
183976      20.0      20.0
183977      7.0      15.0

```

[183978 rows x 42 columns]


In [5]: df.head()

Out[5]:

id	player_fifa	ani_id	player_ani_id	date	overall_rating	potential	preferred
----	-------------	--------	---------------	------	----------------	-----------	-----------

	id	player_id	team_id	date	overall_rating	potential	preferred
0	1	218353	505942	2016-02-18 00:00:00	67.0	71.0	right
1	2	218353	505942	2015-11-19 00:00:00	67.0	71.0	right
2	3	218353	505942	2015-09-21 00:00:00	62.0	66.0	right
3	4	218353	505942	2015-03-20 00:00:00	61.0	65.0	right
4	5	218353	505942	2007-02-22 00:00:00	61.0	65.0	right

5 rows × 42 columns



Creating Target variable:

In [6]: `target = df.pop('overall_rating')`In [7]: `df.shape`

Out[7]: (183978, 41)

In [8]: `target.head()`

```
Out[8]: 0    67.0
1    67.0
2    62.0
3    61.0
4    61.0
Name: overall_rating, dtype: float64
```

Imputing target funtion :

In [9]: `target.isnull().values.sum()`

Out[9]: 836

In [10]: `target.describe()`

```
Out[10]: count    183142.000000
mean         68.600015
std           7.041139
min          33.000000
25%          64.000000
...
```


```

50%      69.000000
75%      73.000000
max       94.000000
Name: overall_rating, dtype: float64

```

```
In [11]: plt.hist(target, 30, range=(33, 94))
```

```
<IPython.core.display.Javascript object>
```



```

C:\ProgramData\Anaconda3\lib\site-packages\numpy\lib\histograms.py:839:
RuntimeWarning: invalid value encountered in greater_equal
    keep = (tmp_a >= first_edge)
C:\ProgramData\Anaconda3\lib\site-packages\numpy\lib\histograms.py:840:
RuntimeWarning: invalid value encountered in less_equal
    keep &= (tmp_a <= last_edge)

```

```

Out[11]: (array([7.0000e+00, 6.0000e+00, 2.0000e+01, 6.5000e+01, 9.4000e+01,
1.4200e+02, 2.9400e+02, 5.2600e+02, 1.2510e+03, 1.8450e+03,
2.5780e+03, 4.0870e+03, 6.1890e+03, 9.6500e+03, 1.3745e+04,
1.8366e+04, 2.0310e+04, 2.0773e+04, 1.9382e+04, 1.8784e+04,
1.5915e+04, 1.1254e+04, 7.5250e+03, 4.9470e+03, 2.8290e+03,
1.4590e+03, 7.4800e+02, 2.2800e+02, 8.4000e+01, 3.9000e+01]),
array([33.          , 35.03333333, 37.06666667, 39.1          , 41.13333333
3,
43.16666667, 45.2          , 47.23333333, 49.26666667, 51.3
,
53.33333333, 55.36666667, 57.4          , 59.43333333, 61.4666666
7,
63.5          , 65.53333333, 67.56666667, 69.6          , 71.63333333
3,
73.66666667, 75.7          , 77.73333333, 79.76666667, 81.8
,
83.83333333, 85.86666667, 87.9          , 89.93333333, 91.9666666
7

```

```

',
    94.    ]),
    <a list of 30 Patch objects>)

```

It's almost normal distribution so we can impute mean value for missing value in target.

```
In [12]: y = target.fillna(target.mean())
```

```
In [13]: y.isnull().values.any()
```

```
Out[13]: False
```

Data Exploration :

```
In [14]: df.columns
```

```
Out[14]: Index(['id', 'player_fifa_api_id', 'player_api_id', 'date', 'potential',
               'preferred_foot', 'attacking_work_rate', 'defensive_work_rate',
               'crossing', 'finishing', 'heading_accuracy', 'short_passing', 'volleys',
               'dribbling', 'curve', 'free_kick_accuracy', 'long_passing',
               'ball_control', 'acceleration', 'sprint_speed', 'agility', 'reactions',
               'balance', 'shot_power', 'jumping', 'stamina', 'strength', 'long_shots',
               'aggression', 'interceptions', 'positioning', 'vision', 'penalties',
               'marking', 'standing_tackle', 'sliding_tackle', 'gk_diving',
               'gk_handling', 'gk_kicking', 'gk_positioning', 'gk_reflexes'],
              dtype='object')
```

```
In [15]: for col in df.columns:
          unique_cat = len(df[col].unique())
          print("{col}--> {unique_cat}..{typ}".format(col=col, unique_cat=unique_cat, typ=df[col].dtype))
```

```

id--> 183978..int64
player_fifa_api_id--> 11062..int64
player_api_id--> 11060..int64
date--> 197..object
potential--> 57..float64
preferred_foot--> 3..object
attacking_work_rate--> 9..object
defensive_work_rate--> 20..object
crossing--> 96..float64
finishing--> 98..float64
heading_accuracy--> 97..float64
short_passing--> 96..float64
volleys--> 94..float64
dribbling--> 98..float64
curve--> 93..float64
free_kick_accuracy--> 98..float64
long_passing--> 96..float64
ball_control--> 94..float64
acceleration--> 87..float64

```


```
sprint_speed--> 86..float64
agility--> 82..float64
reactions--> 79..float64
balance--> 82..float64
shot_power--> 97..float64
jumping--> 80..float64
stamina--> 85..float64
strength--> 83..float64
long_shots--> 97..float64
aggression--> 92..float64
interceptions--> 97..float64
positioning--> 96..float64
vision--> 98..float64
penalties--> 95..float64
marking--> 96..float64
standing_tackle--> 96..float64
sliding_tackle--> 95..float64
gk_diving--> 94..float64
gk_handling--> 91..float64
gk_kicking--> 98..float64
gk_positioning--> 95..float64
gk_reflexes--> 93..float64
```

In [16]: `dummy_df = pd.get_dummies(df, columns=['preferred_foot', 'attacking_work_rate', 'defensive_work_rate'])`
`dummy_df.head()`

Out[16]:

	id	player_fifa_api_id	player_api_id	date	potential	crossing	finishing	heading
0	1	218353	505942	2016-02-18 00:00:00	71.0	49.0	44.0	71.0
1	2	218353	505942	2015-11-19 00:00:00	71.0	49.0	44.0	71.0
2	3	218353	505942	2015-09-21 00:00:00	66.0	49.0	44.0	71.0
3	4	218353	505942	2015-03-20 00:00:00	65.0	48.0	43.0	70.0
4	5	218353	505942	2007-02-22 00:00:00	65.0	48.0	43.0	70.0

5 rows × 67 columns



In [17]: `X = dummy_df.drop(['id', 'date'], axis=1)`

Feature selection :

As tree model doesn't gets affected by missing values present in data set. but feature selection by SelectFromModel can not be done on datasets that carries null value. Therefore, we should also perform imputation on dataset.

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
In [19]: #imputing null value of each column with the mean of that column
imput = Imputer()
X_train = imput.fit_transform(X_train)
X_test = imput.fit_transform(X_test)
```

C:\Users\Shridhar M\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\deprecation.py:66: DeprecationWarning: Class Imputer is deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer from sklearn instead.
warnings.warn(msg, category=DeprecationWarning)

```
In [20]: #finding feature_importance for feature selection. from it we'll be able to decide threshold value
model = XGBRegressor()
model.fit(X_train, y_train)
print(model.feature_importances_)
```

[13:50:34] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```
[0.05115651 0.02897777 0.13911739 0.02826106 0.00595768 0.04332646
 0.01590659 0.          0.01287898 0.          0.          0.00348626
 0.13141693 0.00436716 0.0058711  0.0065089  0.16110069 0.
 0.05513685 0.0040621  0.00863542 0.03761797 0.00898127 0.0139806
 0.01894557 0.02671051 0.0015188  0.          0.03219567 0.0530614
 0.00176505 0.02725989 0.02712863 0.01008714 0.02617038 0.00840923
 0.          0.          0.          0.          0.          0.
 0.          0.          0.          0.          0.          0.
 0.          0.          0.          0.          0.          0.
 0.          0.          0.          0.          0.          0.]
```

```
In [21]: selection = SelectFromModel(model, threshold=0.01, prefit=True)

select_X_train = selection.transform(X_train)
select_X_test = selection.transform(X_test)
select_X_train.shape
```

```
Out[21]: (137983, 20)
```

Scaling the data:

```
In [22]: scalar = StandardScaler()
x_scaled_train = scalar.fit_transform(select_X_train)
x_scaled_train
```

```
Out[22]: array([[ 1.0567811 ,  2.90118168, -0.37370531, ..., -0.69862488,
                  -0.65367807, -0.31949444],
                  [ 0.83239093,  1.11023832, -0.67788964, ..., -0.25617622,
```

```
-0.51352154, -0.25716519],
[ 0.17077907,  1.07420333,  0.38675551, ..., -0.12976231,
-0.60695922, -0.19483593],
...,
[-2.07758255, -0.8212941 ,  1.2993085 , ...,  0.31268635,
 2.4764844 ,  0.30379811],
[ 0.44157109, -0.11639067,  0.99512417, ...,  0.37589331,
 1.02820027,  0.36612736],
[ 0.22002412, -0.64891505,  1.755585 , ..., -0.69862488,
-0.56024038, -0.50648221]]])
```

```
In [23]: x_scaled_test = scalar.fit_transform(select_X_test)
x_scaled_test
```

```
Out[23]: array([[ 0.5825465 ,  0.37577743,  1.14364038, ..., -0.69951284,
-0.65497804, -0.19548251],
[ 0.7131656 ,  0.14588373, -0.22328168, ..., -0.63646818,
-0.60825467, -0.19548251],
[ 0.2307514 , -0.78242149, -0.52704214, ..., -0.51037885,
-0.23446774, -0.00935998],
...,
[ 0.7233016 ,  0.76992259,  0.83987993, ..., -0.13211088,
-0.37463784, -0.13344167],
[ 0.71524098,  1.61776501,  0.6879997 , ..., -0.25820021,
-0.65497804, -0.50568674],
[ 0.90691696,  2.24229255, -1.74208398, ..., -0.25820021,
-0.5615313 , -0.38160505]])
```

Training different models :

1. Linear Regression :

```
In [24]: linear_reg = LinearRegression()
linear_reg.fit(x_scaled_train, y_train)
```

```
Out[24]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normaliz
e=False)
```

```
In [25]: linear_reg.score(x_scaled_test, y_test)
```

```
Out[25]: 0.8536634604512405
```

Hyperparameter Tuning:

```
In [26]: cv = ShuffleSplit(random_state=0)    #defining type of cross_validation
(shuffle splitting)

param_grid = {'n_jobs': [-1]}    #parameters for model tuning

grid = GridSearchCV(linear_reg, param_grid=param_grid, cv=cv)
```

```
In [27]: grid.fit(select_X_train, y_train)
```

```
Out[27]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random state=0, test size=Non
```



```
e, train_size=None),
    error_score='raise-deprecating',
    estimator=LinearRegression(copy_X=True, fit_intercept=True,
                                n_jobs=None, normalize=False),
    iid='warn', n_jobs=None, param_grid={'n_jobs': [-1]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)
```

In [28]: `grid.best_params_`

Out[28]: `{'n_jobs': -1}`

In [29]: `grid.best_estimator_`

Out[29]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1, normalize=False)`

In [30]: `new_linear_reg = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1, normalize=False)`
`new_linear_reg.fit(x_scaled_train, y_train)`

Out[30]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=-1, normalize=False)`

In [31]: `new_linear_reg.score(x_scaled_test, y_test)`

Out[31]: `0.8536634604512405`

2. Decision Tree :

In [32]: `decision_tree = DecisionTreeRegressor(criterion='mse', random_state=0)`
`#estimator`
`decision_tree.fit(x_scaled_train, y_train)`

Out[32]: `DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,`
`max_leaf_nodes=None, min_impurity_decrease=0.0,`
`min_impurity_split=None, min_samples_leaf=1,`
`min_samples_split=2, min_weight_fraction_leaf=0.`
`0,`
`presort=False, random_state=0, splitter='best')`

In [33]: `decision_tree.score(x_scaled_test, y_test)`

Out[33]: `0.9576184490579795`

In [34]: `cv = ShuffleSplit(n_splits=10, random_state=42) #cross validation`
`n`
`param_grid = {'max_depth': [3, 5, 7, 9, 13],`
`'criterion': ['mse', 'friedman_mse']}`
`grid = GridSearchCV(decision_tree, param_grid=param_grid, cv=cv)`

```
In [35]: grid.fit(select_X_train, y_train) #training
```

```
Out[35]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=42, test_size=None, train_size=None),
                    error_score='raise-deprecating',
                    estimator=DecisionTreeRegressor(criterion='mse', max_depth=None,
                                                    max_features=None,
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    presort=False, random_state=0,
                                                    splitter='best'),
                    iid='warn', n_jobs=None,
                    param_grid={'criterion': ['mse', 'friedman_mse'],
                                'max_depth': [3, 5, 7, 9, 13]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring=None, verbose=0)
```

```
In [36]: grid.best_params_
```

```
Out[36]: {'criterion': 'friedman_mse', 'max_depth': 13}
```

```
In [37]: grid.best_estimator_
```

```
Out[37]: DecisionTreeRegressor(criterion='friedman_mse', max_depth=13, max_features=None,
                               max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=1,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               presort=False, random_state=0, splitter='best')
```

```
In [38]: new_deci_tree = DecisionTreeRegressor(criterion='friedman_mse', max_depth=24,
                                                max_features=None, max_leaf_nodes=None,
                                                min_impurity_decrease=0.0, min_impurity_split=None,
                                                min_samples_leaf=1, min_samples_split=2,
                                                min_weight_fraction_leaf=0.0, presort=False, random_state=45,
                                                splitter='best')
new_deci_tree.fit(x_scaled_train, y_train)
```

```
Out[38]: DecisionTreeRegressor(criterion='friedman_mse', max_depth=24, max_features=None,
                               max_leaf_nodes=None, min_impurity_decrease=0.0,
                               min_impurity_split=None, min_samples_leaf=1,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               presort=False, random_state=45, splitter='best')
```

```
In [39]: new_deci_tree.score(x_scaled_test, y_test)
```

```
In [39]: new_model.score(x_scaled_test, y_test)
```

```
Out[39]: 0.9572660408433513
```

3. Random Forest :

```
In [40]: rand_forest = RandomForestRegressor(random_state=123)
rand_forest.fit(x_scaled_train, y_train)
```

C:\Users\Shridhar M\AppData\Roaming\Python\Python37\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

```
Out[40]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10,
                                n_jobs=None, oob_score=False, random_state=123,
                                verbose=0, warm_start=False)
```

```
In [41]: rand_forest.score(x_scaled_test, y_test)
```

```
Out[41]: 0.976235810489465
```

Hyperparameter Tuning:

```
In [42]: cv = ShuffleSplit(test_size=0.2, random_state=0)

param_grid = {'max_features':['sqrt', 'log2', 10],
              'max_depth':[9, 11, 13]}

grid = GridSearchCV(rand_forest, param_grid=param_grid, cv=cv)
```

```
In [43]: grid.fit(x_scaled_train, y_train)
```

```
Out[43]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=0, test_size=0.2,
                                     train_size=None),
                      error_score='raise-deprecating',
                      estimator=RandomForestRegressor(bootstrap=True, criterion=
                      'mse',
                                                         max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=
0.0,
                                                         n_estimators=10, n_jobs=No
                                                         oob_score=False, random st
```

```

ate=123,
                                oob_score=False, random_state=
                                verbose=0, warm_start=False),
                                iid='warn', n_jobs=None,
                                param_grid={'max_depth': [9, 11, 13],
                                'max_features': ['sqrt', 'log2', 10]},
                                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                                scoring=None, verbose=0)

```

In [44]: `grid.best_estimator_`

```

Out[44]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=13,
                                max_features=10, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10,
                                n_jobs=None, oob_score=False, random_state=123,
                                verbose=0, warm_start=False)

```

```

In [45]: new_rand_forest = RandomForestRegressor(bootstrap=True, criterion='mse'
                                , max_depth=33,
                                max_features=10, max_leaf_nodes=None, min_impurity_decrease=
                                0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=10, n_jobs=1, oob_score=False, random_state=42,
                                verbose=0, warm_start=False)
                                new_rand_forest.fit(x_scaled_train, y_train)

```

```

Out[45]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=33,
                                max_features=10, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_
                                jobs=1,
                                oob_score=False, random_state=42, verbose=0,
                                warm_start=False)

```

In [46]: `new_rand_forest.score(x_scaled_test, y_test)`

Out[46]: 0.9790005019563306

4. Xgboost regressor :

```

In [47]: xgr = XGBRegressor(random_state=42)
                                xgr.fit(x_scaled_train, y_train)

```

[14:28:50] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```

Out[47]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                                colsample_bynode=1, colsample_bytree=1, gamma=0,
                                importance_type='gain', learning_rate=0.1, max_delta_step=

```

```
0,  
max_depth=3, min_child_weight=1, missing=None, n_estimator  
s=100,  
n_jobs=1, nthread=None, objective='reg:linear', random_sta
```